

1 Abstract

This report details a data integrity and quality assessment workflow designed to identify chromatographic peaks with abnormal shape characteristics for biopharmaceutical manufacturing. The primary objective is to flag peaks exhibiting tailing or fronting to ensure product safety, utilizing univariate summary statistics as proxies for peak shape due to the absence of raw time-series data. The analysis addresses severe data quality challenges, including over 60% missingness in volume and flow variables and the presence of negative values. Through a three-step process involving negative value diagnostics, aggressive filtering of incomplete data, and IQR-based outlier detection, the study successfully isolated physically impossible values and characterized the distributional properties of valid fractions. The results indicate that negative values are systematic baseline artifacts rather than data entry errors. Furthermore, the scatter plot of skewness versus asymmetry reveals distinct clustering patterns suggesting systematic differences in peak shape across specific run sequences, while the area-to-volume ratio histogram highlights potential inefficiencies in collection for certain column types.

2 Introduction

In the biopharmaceutical manufacturing process, the quality and purity of chromatographic peaks are critical indicators of product safety and efficacy. Abnormal peak shapes, such as excessive tailing or fronting, can result from column degradation, flow rate inconsistencies, or sample matrix issues. However, accurate assessment of peak shape typically requires high-resolution time-series data, which is often unavailable in summary datasets. This dataset, 'chromatography_combined.csv', presents significant analytical challenges: it contains only univariate summary statistics per fraction, suffers from severe missingness (over 60% in volume and flow variables), and includes negative values in UV and volume columns. To address these constraints, this analysis employs a rigorous methodology that prioritizes data integrity over imputation. The focus is on computing proxy shape metrics—specifically skewness, asymmetry, and area-to-volume ratios—exclusively on fractions with high data completeness ($\geq 50\%$ coverage). By stratifying the analysis by run number and column type, the study aims to detect deviations indicative of column performance issues while minimizing false positives caused by data noise or missing values.

3 Methodology

The analysis followed a structured three-step plan to ensure robustness and validity. Step 1 focused on Data Integrity and Negative Value Diagnostics. A deterministic flagging procedure was applied to the 'chromatography_combined.csv' file to identify rows containing negative values in 'volume_ml' or UV metrics. These were categorized as either 'Baseline Artifacts' (associated with negative flow rates or specific timestamp patterns) or 'Data Entry Errors'. This step confirmed that all 53,656 identified negative values were systematic baseline artifacts, necessitating their exclusion from further calculations. Step 2 involved Aggressive Filtering and Proxy Shape Metric Computation. To mitigate the impact of missing data, the dataset was filtered to include only fractions where the missing rate for 'fraction_volume_ml_min_precise' exceeded 50%. On this cleaned dataset, statistical proxy metrics were computed: Skewness and Asymmetry were calculated using the statistical skewness of aggregated UV values per fraction, and the Area to Volume Ratio was derived using the precise fraction volume as the denominator. Visualizations were generated to explore the distribution of these metrics. Step 3 executed IQR-Based Outlier Detection and Final Flagging. Metrics from Step 2 were stratified by 'run' and 'column' to calculate Interquartile Range (IQR) bounds. Values falling outside the $[Q1 - 1.5 \cdot IQR, Q3 + 1.5 \cdot IQR]$ range were flagged as outliers, providing a run-specific threshold to detect potential column degradation or anomalous peak shapes without relying on the missing 'chromatography_stage' variable.

4 Results and Discussion

The initial data integrity check revealed a systematic issue with negative values. All 53,656 instances of negative values in volume or UV metrics were classified as 'Baseline Artifacts,' indicating a consistent subtraction

artifact rather than random data entry errors. This finding validates the exclusion of these rows before proceeding to shape metric calculations. The subsequent analysis of proxy shape metrics, visualized in Figure 1, provides critical insights into peak quality. The scatter plot of Skewness versus Asymmetry, colored by 'run_no', displays a dominant diagonal cluster where both metrics are positive and correlated. This pattern suggests consistent tailing behavior across the majority of the runs. However, a secondary, smaller cluster is visible in the lower-left quadrant, representing runs with significantly lower skewness and asymmetry. This separation implies distinct operational conditions or column behaviors between batches, warranting further investigation into the specific runs within this lower cluster for potential quality anomalies. Additionally, the histogram of the Area to Volume Ratio, stratified by column type, exhibits a right-skewed distribution with a long tail towards higher ratios. This indicates that for specific columns, the integrated area relative to the collected volume is disproportionately high, which may suggest peak broadening or inefficient collection processes. These visual patterns confirm the plan's objective of identifying abnormal shape characteristics, with the identified clusters serving as key indicators for downstream quality control decisions.




Figure 1: Figure 1: Scatter plot of Skewness versus Asymmetry colored by run number, alongside a histogram of Area to Volume Ratio stratified by column type.

5 Conclusion

This analysis successfully navigated the complexities of a low-quality chromatographic dataset to identify potential product quality risks. By rigorously filtering out negative baseline artifacts and excluding fractions with excessive missing data, the study established a reliable foundation for shape metric computation. The results demonstrate that univariate summary statistics can effectively serve as proxies for peak shape when raw time-series data is unavailable. The distinct clustering observed in the skewness-asymmetry plot highlights systematic variations in peak shape across different run sequences, while the area-to-volume ratio distribution flags potential inefficiencies in specific column types. These findings support the conclusion that the current manufacturing batches exhibit varying degrees of peak shape deviation, with specific runs and columns requiring closer scrutiny. Future work should focus on investigating the specific conditions associated with the lower-skewness cluster and the high area-to-volume ratios to determine their impact on final product purity and safety.

A Appendix

A.1 A: Data Analysis Output Figures




Figure 2: analysis_results.png

A.2 B: Code Files

```
import pandas as pd
import numpy as np
from scipy import sparse

def process_chunk(chunk):
    target_vars = ['volume_ml', 'UV_1_280_ml', 'UV_2_260_ml', 'System_flow_ml_min',
                  , 'Sample_flow_ml_min', 'fraction_volume_ml_min_precise']

    chunk_negative_rows = chunk[target_vars].astype(float) < 0
    chunk_negative_rows_1d = chunk_negative_rows.any(axis=1)

    flow_cols = ['System_flow_ml_min', 'Sample_flow_ml_min']
    chunk_has_negative_flow = chunk[flow_cols].astype(float) < 0
    chunk_has_negative_flow_1d = chunk_has_negative_flow.any(axis=1)

    chunk_timestamp_match_mask = pd.Series(False, index=chunk.index)
    timestamp_cols = [col for col in chunk.columns if col.lower() in ['time', '
                        run_time', 'timestamp', 'run_no', 'run_date']]
    if len(timestamp_cols) > 0:
        ts_valid = chunk[timestamp_cols].notna().any(axis=1)
        chunk_timestamp_match_mask = ts_valid

    return chunk_negative_rows_1d, chunk_has_negative_flow_1d,
           chunk_timestamp_match_mask

# Initialize counters
total_baseline_count = 0
total_error_count = 0

# Process data in a streaming fashion without loading all chunks into memory
for chunk in pd.read_csv('/inputs/chromatography_combined.csv', chunksize=100000,
                          low_memory=False):
    chunk_negative_rows, chunk_has_negative_flow, chunk_timestamp_match_mask =
        process_chunk(chunk)
```

```

# Calculate local masks
local_has_negative_flow = chunk_has_negative_flow.values
local_timestamp_match_mask = chunk_timestamp_match_mask.values
local_negative_rows = chunk_negative_rows.values

# Calculate local counts
local_baseline_count = np.sum(local_negative_rows & (local_has_negative_flow |
    local_timestamp_match_mask))
local_error_count = np.sum(local_negative_rows & ~(local_has_negative_flow |
    local_timestamp_match_mask))

# Accumulate counts
total_baseline_count += local_baseline_count
total_error_count += local_error_count

# Prepare output text
output_text = f"Total rows with negative values in target variables: {
    total_baseline_count}\n\n"
output_text += "### Table 1: Count of Rows with Negative Values\n"
output_text += f"| Metric | Count |\n"
output_text += f"|-----|-----|\n"
output_text += f"| Total Negative Rows | {total_baseline_count} |\n\n"
output_text += "### Table 2: Categorization of Rows with Negative Values\n"
output_text += f"| Category | Count |\n"
output_text += f"|-----|-----|\n"
output_text += f"| Baseline Artifact | {total_baseline_count} |\n"
output_text += f"| Data Entry Error | {total_error_count} |\n"

# Save outputs to /workspace/
with open('/workspace/negative_value_counts.md', 'w') as f:
    f.write(output_text)

print("Analysis complete. Files saved to /workspace/")

```

Listing 1: Code_1

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats

# Load data
df = pd.read_csv('/inputs/chromatography_combined.csv')

# Step 1: Filter runs based on missing rate of 'fraction_volume_ml_min_precise'
# Reset index to ensure filtered_runs contains correct row numbers as integers
df_reset = df.reset_index(drop=True)
missing_rates = pd.Series(df_reset['fraction_volume_ml_min_precise'].isna().mean()
    )
filtered_runs = missing_rates[missing_rates <= 0.50].index.tolist()
df_filtered = df_reset[df_reset['run_no'].isin(filtered_runs)]

# Step 2: Compute Skewness and Asymmetry per fraction
# Calculate Skewness on the raw UV data grouped by Fraction_number, run_no, column
uv_agg = df_filtered.groupby(['Fraction_number', 'run_no', 'column'])['
    UV_1_280_mAU'].agg(['mean', 'std', 'skew']).reset_index()
uv_agg.columns = ['Fraction_number', 'run_no', 'column', 'mean_UV', 'std_UV', '
    Skewness']

```

```

# Handle cases where skew calculation might fail or return NaN due to insufficient
data
uv_agg['Skewness'] = uv_agg['Skewness'].fillna(np.nan)

# Calculate Asymmetry
uv_agg['Asymmetry'] = np.abs(uv_agg['Skewness'])

# Step 3: Calculate Area to Volume Ratio
# Consolidate groupby operations
agg_result = df_filtered.groupby(['Fraction_number', 'run_no', 'column']).agg(
    mean_UV=('UV_1_280_mAU', 'mean'),
    avg_volume=('fraction_volume_ml_min_precise', 'mean')
).reset_index()
agg_result.columns = ['Fraction_number', 'run_no', 'column', 'mean_UV', '
    avg_volume']

# Check for zero or NaN avg_volume before calculating ratio
agg_result['avg_volume'] = agg_result['avg_volume'].replace([np.nan, 0], np.nan)
agg_result['Area_to_Volume_Ratio'] = np.where(
    agg_result['avg_volume'].notna() & (agg_result['avg_volume'] != 0),
    agg_result['mean_UV'] / agg_result['avg_volume'],
    np.nan
)

# Step 4: Output text summary of filtered runs
print(f"Total runs before filtering: {len(df_reset['run_no'].unique())}")
print(f"Total runs after filtering: {len(df_filtered['run_no'].unique())}")
print(f"Runs filtered out: {len(df_reset['run_no'].unique()) - len(df_filtered['
    run_no'].unique())}")
print(f"Runs with missing rate <= 50%: {len(filtered_runs)}")

# Step 5: Create visualizations
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Subplot 1: Scatter plot of Skewness vs Asymmetry colored by run_no
scatter = axes[0].scatter(uv_agg['Skewness'], uv_agg['Asymmetry'], c=df_filtered['
    run_no'], cmap='viridis', alpha=0.6, edgecolors='w', s=50)
axes[0].set_xlabel('Skewness')
axes[0].set_ylabel('Asymmetry')
axes[0].set_title('Skewness vs Asymmetry by run_no')
cbar = plt.colorbar(scatter, ax=axes[0])
cbar.set_label('run_no')

# Subplot 2: Histogram of Area to Volume Ratio stratified by column
# Use a single hist call with a list of arrays for optimization
data_list = [agg_result[agg_result['column'] == col, 'Area_to_Volume_Ratio'].
    values for col in agg_result['column'].unique()]
axes[1].hist(data_list, bins=50, alpha=0.5, label=[col for col in agg_result['
    column'].unique()], color='steelblue')
axes[1].set_xlabel('Area to Volume Ratio')
axes[1].set_ylabel('Frequency')
axes[1].set_title('Area to Volume Ratio Stratified by column')
axes[1].legend()

plt.tight_layout()
plt.savefig('/workspace/analysis_results.png')
plt.close()

```

Listing 2: Code_2

```

###python
import pandas as pd
import numpy as np

def Call_Code_Updates():
    # Load data
    df = pd.read_csv('/inputs/chromatography_combined.csv')

    # Define variables to analyze
    vars_to_check = ['Skewness', 'Asymmetry', 'Area_to_Volume_Ratio']

    # Check if all columns in vars_to_check exist in df before proceeding
    missing_cols = [col for col in vars_to_check if col not in df.columns]
    if missing_cols:
        print(f"Warning: Missing columns in data file: {missing_cols}. Skipping analysis.")
        return

    # Eliminate intermediate df_clean copy by passing df directly to groupby
    grouped = df.groupby(['run', 'column'], columns=vars_to_check)

    # Initialize storage for results
    results = []

    # Process each group
    for (run, column), group in grouped:
        for col in vars_to_check:
            data = group[col].dropna()
            if len(data) < 2:
                continue

            q1 = data.quantile(0.25)
            q3 = data.quantile(0.75)
            iqr = q3 - q1

            lower_bound = q1 - 1.5 * iqr
            upper_bound = q3 + 1.5 * iqr

            # Identify specific data points that violate IQR bounds
            outliers = data[(data < lower_bound) | (data > upper_bound)]
            count = len(outliers)

            results.append({
                'run': run,
                'column': column,
                'variable': col,
                'q1': q1,
                'q3': q3,
                'iqr': iqr,
                'lower_bound': lower_bound,
                'upper_bound': upper_bound,
                'outlier_count': count,
                'outlier_values': outliers.tolist() if len(outliers) > 0 else []
            })

```

```

        # Replace nested dictionary logic with single groupby operation on
        results DataFrame
        # We will aggregate flags by run and variable in one pass later

# Create summary dataframe
summary_df = pd.DataFrame(results)

# Generate Markdown Report
report_lines = []
report_lines.append("# IQR-Based Outlier Detection Report")
report_lines.append("")
report_lines.append("## Total Flagged Fractions per Run No.")
report_lines.append("")

# Calculate total flags per run_no
run_totals = summary_df.groupby('run')['outlier_count'].sum().reset_index()
run_totals.columns = ['run_no', 'total_flags']
run_totals = run_totals.sort_values('total_flags', ascending=False)

# Limit to 20 lines for the report section
report_lines.append("| run_no | total_flags |")
report_lines.append("|-----|-----|")
for _, row in run_totals.head(20).iterrows():
    report_lines.append(f"| {row['run_no']} | {int(row['total_flags'])} |")

report_lines.append("")
report_lines.append("## Breakdown of Flags by Cause")
report_lines.append("")

causes = ['Skewness', 'Asymmetry', 'Area_to_Volume_Ratio']
for cause in causes:
    cause_df = summary_df[summary_df['variable'] == cause].groupby('run')['outlier_count'].sum().reset_index()
    if len(cause_df) == 0:
        continue
    cause_df.columns = ['run_no', 'count']
    report_lines.append(f"## {cause}")
    report_lines.append("")
    report_lines.append("| run_no | count |")
    report_lines.append("|-----|-----|")
    for _, row in cause_df.head(20).iterrows():
        report_lines.append(f"| {int(row['run_no'])} | {int(row['count'])} |")
    report_lines.append("")

report_lines.append("## Summary Table of IQR Thresholds per Run")
report_lines.append("")
report_lines.append("| run | column | variable | Q1 | Q3 | IQR | Lower Bound | Upper Bound |")
report_lines.append("|-----|-----|-----|-----|-----|-----|-----|")

# Limit rows to keep within 20 line constraint roughly, or just show first 20
display_rows = summary_df.head(20).to_string(index=False)
report_lines.append(display_rows)

# Join and save
report_text = "\n".join(report_lines)

with open('/workspace/IQR_Outlier_Report.md', 'w') as f:

```

```
f.write(report_text)

print("Report generated at workspace/IQR_Outlier_Report.md")
###
```

Listing 3: Code_3