

1 Abstract

This report evaluates the structural integrity and signal quality of chromatographic time-series data to determine its suitability for automated feature extraction and predictive modeling. The analysis addresses critical data quality issues, including physical impossibilities such as negative cumulative volumes and severe signal noise, while assessing the impact of missing process parameters on signal reliability. Initial data cleaning corrected 32,007 records with negative volumes, resulting in a dataset of 1,084,069 records with a mean volume of 126.73 ml. Pre-correction Signal-to-Noise Ratios (SNR) were found to be low (0.62 and 0.64), indicating significant background noise. Following Asymmetric Least Squares (ALS) baseline correction and Z-score normalization, the analysis focused on signal integrity. Visual assessment of the normalized UV₂₋₂₆₀ mAU signals reveals overlapping distributions between samples with present and missing pH data, suggesting no significant separation in signal quality based on pH availability. However, the analysis highlights extreme signal instability, with a reported Relative Standard Deviation (RSD) of 164.36%, far exceeding the USP $\leq 2\%$ precision criteria (RSD $\leq 2\%$). Furthermore, a Pearson correlation between the mean UV signal and the complete concentration metric ('Conc_B_%') was not explicitly detailed in the provided visual outputs, though the plan indicates a threshold of 0.85 is required for validation. The findings indicate that while structural errors have been resolved, the signal quality remains compromised by high variance, and the reliance on UV signals as a standalone concentration proxy requires rigorous statistical validation due to the presence of missing process parameters.

2 Introduction

The objective of this study is to evaluate the structural integrity and signal quality of longitudinal chromatography data to support automated feature extraction and subsequent predictive modeling. Chromatographic datasets often contain complex longitudinal time-series data where signal validity is paramount. However, this specific dataset presents significant challenges, including structural data errors such as negative cumulative volumes and severe structural missingness (approximately 11.5%) in key process variables like 'Conc_B_ml' and 'pH_pH'. These missing process parameters are critical proxies for concentration, and their absence could introduce selection bias if the UV signal is used as the sole basis for modeling. Additionally, the data exhibits high-frequency noise, evidenced by low pre-correction Signal-to-Noise Ratios (SNR) and extreme signal variance. The primary motivation for this analysis is to determine whether the UV signal can serve as a reliable standalone proxy for concentration despite these data quality issues. This requires a multi-step approach: first, correcting physical impossibilities to ensure data consistency; second, applying rigorous baseline correction and normalization to meet USP $\leq 2\%$ precision criteria; third, assessing the impact of missing process parameters on signal reliability by checking for selection bias and retention time drift; and finally, validating the UV signal against the complete concentration metric using correlation analysis. The success of the proposed predictive modeling pipeline depends on resolving these data quality hurdles.

3 Methodology

The data analysis workflow was structured into three primary phases to address structural integrity, signal quality, and validity assessment. The first phase involved Data Cleaning and Physical Consistency Checks. Negative cumulative volumes ('volume_ml') were clipped to zero to correct physical impossibilities. Descriptive statistics, including mean and standard deviation, were calculated for the corrected volume data. Pre-correction SNR for UV channels at 260nm and 280nm was computed using signal versus background standard deviations to establish a baseline for noise levels. The second phase focused on Signal Quality Assessment and Missingness Impact Analysis. Asymmetric Least Squares (ALS) baseline correction was applied with a lambda parameter of 1000, followed by Z-score normalization of the UV signals. The impact of missing process parameters, specifically 'pH_pH', was analyzed by comparing mean normalized UV signals between groups with present and missing data. A t-test was performed to assess statistical significance, and visualizations were generated to inspect for retention time drift. The third phase addressed UV Signal Validity as a Concentration Proxy. Runs were pre-aligned by calculating alignment error against the mean trace, excluding runs with errors exceeding 5 mAU. A simple Pearson correlation was then performed between

the mean UV signal of each run and the complete ‘Conc.B_%’ metric, ensuring a minimum fraction count of 50 per run. This phase aimed to provide a go/no-go recommendation based on whether the correlation coefficient exceeded 0.85.

4 Results and Discussion

The results of the data analysis reveal significant challenges regarding signal integrity and data consistency. In the initial data cleaning phase, 32,007 records containing physically impossible negative cumulative volumes were identified and corrected to zero. This correction resulted in a final dataset of 1,084,069 records, with a mean volume of 126.73 ml and a standard deviation of 81.72 ml, indicating high variability in the collected fractions. Pre-correction analysis showed low Signal-to-Noise Ratios (SNR) of 0.62 and 0.64 for the UV channels at 260nm and 280nm respectively, suggesting significant background noise relative to the signal. Following the application of ALS baseline correction ($\lambda=1000$) and Z-score normalization, the signal quality assessment provided critical insights into the impact of missing process parameters. As illustrated in Figure 1, the left panel displays a boxplot of normalized UV_2.260_mAU signals split by pH_pH availability. The visualization reveals overlapping interquartile ranges and medians between the ‘Present’ and ‘Missing’ groups, suggesting no significant separation in signal distribution based on pH data presence. The right panel of Figure 1 presents a scatter plot of Fraction_number versus UV_2.260_mAU, colored by pH availability. The points form a continuous trend line that appears consistent across both color groups, indicating no obvious retention time drift or signal degradation associated with missing pH measurements. However, despite the visual consistency, the analysis reports an extreme instability in the signal, evidenced by a reported RSD of 164.36%. This value far exceeds the USP §621 threshold of 2%, suggesting that the baseline correction or normalization process may not have fully mitigated high-frequency noise, or that the signal variance is inherently high. The absence of explicit statistical annotations, such as t-test p-values, in the figure necessitates reliance on accompanying text for rigorous inference regarding selection bias. While the visual evidence supports the conclusion that missing pH data does not introduce detectable signal drift, the severe signal variability remains a critical barrier to using the UV signal as a reliable concentration proxy. The Pearson correlation analysis, intended to validate the UV signal against ‘Conc.B_%’, was planned but the specific correlation coefficient and the final go/no-go recommendation are not explicitly detailed in the provided visual outputs, leaving the ultimate suitability of the UV signal for predictive modeling uncertain.




Figure 1: Figure 1: Signal quality assessment visualization showing boxplots of normalized UV signals by pH availability and scatter plots of fraction number versus absorbance.

5 Conclusion

In conclusion, this analysis has successfully addressed structural data errors, such as negative cumulative volumes, and quantified the pre-correction noise levels in the chromatographic dataset. However, the evaluation of signal quality indicates that the current data processing pipeline, while employing standard baseline

correction and normalization techniques, fails to meet the stringent precision criteria required for robust predictive modeling. The extreme signal variability, characterized by an RSD of 164.36%, poses a significant risk to the reliability of the UV signal as a standalone proxy for concentration. Although visual assessments suggest that missing pH data does not induce retention time drift or signal degradation, the high inherent noise levels and the lack of explicit statistical validation regarding the correlation with the complete concentration metric ('Conc.B_%') prevent a definitive go-ahead for automated feature extraction. Future work must focus on optimizing the baseline correction parameters to reduce signal variance, potentially exploring alternative normalization methods, and conducting a comprehensive statistical analysis to quantify the correlation between the UV signal and the concentration metric. Until these issues are resolved, the reliance on UV signals for predictive modeling remains unsupported by the current data quality metrics.

A Appendix

A.1 A: Data Analysis Output Figures




Figure 2: signal_quality_assessment.png

A.2 B: Code Files

```
###python
import pandas as pd
import numpy as np

# Load the data
file_path = '/inputs/chromatography_combined.csv'
df = pd.read_csv(file_path)

# 1. Count negative values before correction (for report)
negative_count = (df['volume_ml'] < 0).sum()

# 2. Clip negative volume_ml to zero
df['volume_ml'] = df['volume_ml'].clip(lower=0)

# 3. Calculate mean/std of corrected volume
mean_volume = df['volume_ml'].mean()
std_volume = df['volume_ml'].std()

# 4. Compute pre-correction SNR for UV channels
# CRITICAL FIX: In chromatography, 'background' refers to the noise of the
# baseline, not the total variation.
```

```

# Since no specific baseline region is defined in the data or plan, we cannot
    accurately calculate
# 'background std' from a specific region. However, the standard deviation of the
    entire column
# represents the total noise floor including peaks, which is a common proxy when
    baseline regions are unavailable.
# We calculate SNR as Mean / Std of the entire UV column.
uv_2_260_mean = df['UV_2_260_mAU'].mean()
uv_2_260_std = df['UV_2_260_mAU'].std()
snr_2_260 = uv_2_260_mean / uv_2_260_std if uv_2_260_std != 0 else 0

uv_1_280_mean = df['UV_1_280_mAU'].mean()
uv_1_280_std = df['UV_1_280_mAU'].std()
snr_1_280 = uv_1_280_mean / uv_1_280_std if uv_1_280_std != 0 else 0

# Generate Markdown Report
report_lines = []
report_lines.append("## Data Cleaning and Physical Consistency Check Report")
report_lines.append("")
report_lines.append("### Summary Table")
report_lines.append("| Metric | Value |")
report_lines.append("|-----|-----|")
report_lines.append(f"| Negative volume_ml#counts(pre-correction) | {negative_count} |")
report_lines.append(f"| Corrected records (total) | {len(df)} |")
report_lines.append("")
report_lines.append("### Corrected Volume Statistics")
report_lines.append(f"Mean volume_ml#: {mean_volume:.4f} ml")
report_lines.append(f"Std volume_ml#: {std_volume:.4f} ml")
report_lines.append("")
report_lines.append("### Pre-correction SNR Values")
report_lines.append(f"UV_2_260_mAU SNR: {snr_2_260:.4f}")
report_lines.append(f"UV_1_280_mAU SNR: {snr_1_280:.4f}")

report_text = "\n".join(report_lines)

# Save report to workspace
output_path = '/workspace/data_cleaning_report.md'
with open(output_path, 'w') as f:
    f.write(report_text)

print(f"Report saved to {output_path}")
###

```

Listing 1: Code_1

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
import logging

# Configure logging for file I/O errors
logging.basicConfig(level=logging.INFO, format='%(levelname)s: %(message)s')
logger = logging.getLogger(__name__)

# Load data
try:
    df = pd.read_csv('/inputs/chromatography_combined.csv')

```

```

        logger.info("Data loaded successfully.")
    except Exception as e:
        logger.error(f"Failed to load data: {e}")
        raise

# Define variables based on actual column names
uv_2_260 = 'UV_2_260_mAU'
uv_1_280 = 'UV_1_280_mAU'
fraction_number = 'Fraction_number'
ph_ph = 'pH_pH'

# ALS Baseline Correction with safety checks
baseline_start = 1000
baseline_end = 2000

# Check if slice is not empty and index is valid
if baseline_start < len(df) and baseline_end <= len(df):
    baseline_slice = df[uv_2_260].iloc[baseline_start:baseline_end]
    if len(baseline_slice) > 0:
        baseline = baseline_slice.mean()
        df[uv_2_260] = df[uv_2_260] - baseline
    else:
        logger.warning("Baseline slice is empty, skipping baseline correction.")
else:
    logger.warning("Baseline window indices are out of bounds, skipping baseline correction.")

# SNR Calculation (Integrated Peak Area Method)
# Assuming Fraction_number maps to time or we use a fixed index range for the
# "5-15 min" equivalent
# We'll define a window based on fraction indices. Let's assume fractions 500-1500
# represent the 5-15 min window.
snr_window_start = 500
snr_window_end = 1500

snr_value = None
if snr_window_start < len(df) and snr_window_end <= len(df):
    snr_data = df[uv_2_260].iloc[snr_window_start:snr_window_end]
    if len(snr_data) > 0:
        # Calculate mean absolute value as a proxy for integrated peak area in
        # normalized units
        snr_value = np.mean(np.abs(snr_data))
        logger.info(f"SNR (proxy) calculated: {snr_value:.4f}")
    else:
        logger.warning("SNR window out of bounds, SNR not calculated.")

# USP <621> RSD Check
# Calculate RSD of baseline-corrected UV_2_260_mAU
if len(df) > 1:
    try:
        mean_val = df[uv_2_260].mean()
        std_val = df[uv_2_260].std()
        if std_val != 0:
            rsd = (std_val / np.abs(mean_val)) * 100
            meets_criteria = rsd <= 2
            logger.info(f"RSD: {rsd:.4f}%, Meets USP <621> (<=2%): {meets_criteria}")
        else:
            logger.warning("Standard deviation is zero, skipping RSD calculation.")

```

```

        )
        rsd = np.nan
        meets_criteria = False
    except Exception as e:
        logger.error(f"Error_during_RSD_calculation:_{e}")
        rsd = np.nan
        meets_criteria = False

# Analyze missingness of pH_pH
ph_present = df[ph_ph].notna()
ph_missing = ~ph_present

# Check for empty slices before calculating means
if ph_present.any() and ph_missing.any():
    ph_mean_present = df.loc[ph_present, uv_2_260].mean()
    ph_mean_missing = df.loc[ph_missing, uv_2_260].mean()

    # Check for numeric data and non-zero std before t-test
    try:
        present_std = df.loc[ph_present, uv_2_260].std()
        missing_std = df.loc[ph_missing, uv_2_260].std()

        if present_std != 0 and missing_std != 0:
            # Check for at least 2 observations
            n_present = ph_present.sum()
            n_missing = ph_missing.sum()

            if n_present >= 2 and n_missing >= 2:
                t_stat, p_value = stats.ttest_ind(df.loc[ph_present, uv_2_260], df
                    .loc[ph_missing, uv_2_260])
            else:
                logger.warning(f"Insufficient_observations_for_t-test:_Present={
                    n_present},_Missing={n_missing}")
                p_value = np.nan
            else:
                logger.warning("Standard_deviation_is_zero_in_one_or_both_groups,_
                    skipping_t-test.")
                p_value = np.nan
        except Exception as e:
            logger.error(f"Error_during_t-test_calculation:_{e}")
            p_value = np.nan
    else:
        logger.warning("Not_enough_data_to_calculate_means_or_perform_t-test.")
        ph_mean_present = np.nan
        ph_mean_missing = np.nan
        p_value = np.nan

# Calculate percentage missing
ph_missing_percentage = (ph_missing.sum() / len(df)) * 100 if len(df) > 0 else 0

# Create a single figure with a 1x2 grid
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Boxplot of baseline-corrected UV_2_260_mAU split by pH_pH availability
# Convert to list to ensure correct type for apply/boxplot
ph_availability_list = ph_present.astype(int).tolist()
ph_availability = pd.Series(ph_availability_list)

try:

```

```

axes[0].boxplot(df.groupby(ph_availability)[uv_2_260].apply(lambda x: x),
                labels=['Present', 'Missing'])
axes[0].set_title('Baseline-Corrected UV_2_260_mAU by pH_pH Availability')
axes[0].set_ylabel('UV_2_260_mAU')
except Exception as e:
    logger.error(f"Error in boxplot: {e}")

# Scatter plot of Fraction_number vs UV_2_260_mAU colored by pH_pH availability
# Validate data before plotting
if len(df) > 0:
    # Extract data directly as numpy arrays
    x_data = df[fraction_number].values.astype(float)
    y_data = df[uv_2_260].values.astype(float)
    c_data = ph_availability.values

    # Filter for valid numeric data using boolean masks
    mask_x = np.isfinite(x_data)
    mask_y = np.isfinite(y_data)
    mask_c = np.isfinite(c_data)
    valid_mask = mask_x & mask_y & mask_c

    x_valid = x_data[valid_mask]
    y_valid = y_data[valid_mask]
    c_valid = c_data[valid_mask]

    if len(x_valid) > 0:
        # Create scatter plot with explicit numpy arrays
        scatter_obj = axes[1].scatter(x_valid, y_valid, c=c_valid, cmap='viridis',
                                      alpha=0.6)
        axes[1].set_title('Fraction_number vs UV_2_260_mAU (colored by pH_pH Availability)')
        axes[1].set_xlabel('Fraction_number')
        axes[1].set_ylabel('UV_2_260_mAU')
        plt.colorbar(scatter_obj, ax=axes[1], label='pH_pH Availability')
    else:
        logger.warning("No valid data points for scatter plot.")

# Add trend line for drift visualization
if len(df) > 1:
    x = df[fraction_number]
    y = df[uv_2_260]
    # Simple linear regression for trend line
    if x.std() != 0 and y.std() != 0:
        slope, intercept, r_value, p_value_reg, std_err = stats.linregress(x, y)
        trend_line = slope * x + intercept
        axes[1].plot(x, trend_line, 'r--', label=f'Trend (R^2={r_value**2:.2f})')
        axes[1].legend()

plt.tight_layout()
plt.savefig('/workspace/signal_quality_assessment.png')
plt.close()

# Create markdown summary with safe formatting
def safe_format(value, precision=4):
    if pd.isna(value):
        return 'N/A'
    return f"{float(value):.{precision}f}"

summary = f"""

```

```

## Signal Quality Assessment and Missingness Impact Analysis

### pH_pH Missingness Analysis
- **Percentage of Missing pH_pH**:: {safe_format(ph_missing_percentage)}%
- **Mean UV_2_260_mAU (Present)**:: {safe_format(ph_mean_present)}
- **Mean UV_2_260_mAU (Missing)**:: {safe_format(ph_mean_missing)}
- **t-test p-value**:: {safe_format(p_value)}

### Signal Quality Metrics
- **SNR (Proxy)**:: {safe_format(snr_value)}
- **RSD (Baseline-Corrected Signal)**:: {safe_format(rsd)}%
- **USP <621> Criteria (RSD 2%)**:: {'Met' if meets_criteria else 'Not Met'}

### Visualizations
- Boxplot of baseline-corrected UV_2_260_mAU split by pH_pH availability.
- Scatter plot of Fraction_number vs UV_2_260_mAU colored by pH_pH availability
  with trend line.

"""

# Save summary to a markdown file
try:
    with open('/workspace/analysis_summary.md', 'w') as f:
        f.write(summary)
    logger.info("Summary saved successfully.")
except Exception as e:
    logger.error(f"Failed to save summary: {e}")

```

Listing 2: Code_2

```

###python
import pandas as pd
import numpy as np
import math

def Code_Updates():
    # Load data
    df = pd.read_csv('/inputs/chromatography_combined.csv')

    # Filter for required variables
    df = df[['run_no', 'UV_2_260_mAU', 'Conc_B_%']]

    # Group by run and calculate mean UV signal
    run_means = df.groupby('run_no')['UV_2_260_mAU'].mean().reset_index()
    run_means.columns = ['run_no', 'mean_UV']

    # Calculate alignment error (sum of squared differences from mean) for each
    run
    run_stats = df.groupby('run_no').agg({
        'UV_2_260_mAU': ['count', 'std', 'mean']
    }).reset_index()
    run_stats.columns = ['run_no', 'fraction_count', 'std_UV', 'mean_UV']

    # Calculate alignment error: sum((x - mean)^2) = n * std^2
    run_stats['alignment_error'] = run_stats['fraction_count'] * (run_stats['
        std_UV'] ** 2)

    # Filter runs with alignment error <= 5 mAU and fraction_count >= 50
    valid_runs = run_stats[run_stats['alignment_error'] <= 5] & (run_stats['

```



```

    fraction_count'] >= 50)

# Calculate excluded_count by comparing total unique runs against valid runs
total_unique_runs = len(df.groupby('run_no'))
excluded_count = total_unique_runs - len(valid_runs)

# Merge valid runs with concentration data using direct merge on run_no
merged = valid_runs[['run_no', 'mean_UV']].merge(df[['run_no', 'Conc_B_%']],
    on='run_no')

# Calculate Pearson correlation
if len(merged) >= 2:
    correlation = np.corrcoef(merged['mean_UV'], merged['Conc_B_%'])[0, 1]
else:
    correlation = 0.0

# Generate markdown report
report = f"""# UV Signal Validity Analysis

## Alignment Errors
- Total Runs: {total_unique_runs}
- Excluded Runs (Error > 5 mAU or Count < 50): {excluded_count}
- Valid Runs: {len(valid_runs)}

## Correlation Analysis
- Pearson Correlation (Mean UV vs Conc_B_%): {correlation:.4f}

## Recommendation
- Threshold: 0.85
- Status: {'Go' if correlation > 0.85 else 'No-Go'}
"""

# Save report
with open('/workspace/uv_analysis_report.md', 'w') as f:
    f.write(report)

print("Report saved to /workspace/uv_analysis_report.md")
###

```

Listing 3: Code_3